

Hash tables

Goldsmiths Computing

Motivation

A different way to implement a collection, with different performance implications

Definition

A hash table is a data structure that can represent a set, or more generally a map of keys to values (an associative array), by computing a numeric value for each key using a **hash function** and then using that numeric value to compute an index into an array to look up the value.

Set operations

`insert[o]` insert the object `o` into the set

`find[o]` is the object `o` in the set?

and also

`delete[o]` delete the object `o` from the set

Sets of small integers

Represent sets of non-negative integers smaller than N using an array of size N . e.g. for domain $[0,5]$:

✓	✗	✗	✓	✗	✓
---	---	---	---	---	---

represents the set $\{0, 3, 5\}$

insert[o] $S[o] \leftarrow \text{true}$

find[o] **return** $S[o]$

delete[o] $S[o] \leftarrow \text{false}$

Sets of unbounded integers

Apply the same representation?

✓	✗	✗	✓	✗	✓	...									
---	---	---	---	---	---	-----	--	--	--	--	--	--	--	--	--

2^{32} integers? 2^{29} bytes of RAM (512MB)

Sets of unbounded integers

If the expected size of the sets is small (even if the range of possible values is large):

1. choose a reasonable size for the array, say twice expected size
2. reduce the integer to within the range of array indices using a function $f(n)$
3. store the (unreduced) integer in the array slot

Then

insert[o] $S[f(o)] \leftarrow o$

find[o] **return** $S[f(o)] = o$


delete[o] $S[f(o)] \leftarrow \text{NIL}$

Example

range $[0, 2^{10})$

set size 5

Choose array size of (say) 11 and compute index as $f(n) = n \bmod 11$



represents the set $\{155, 167, 603, 951, 1004\}$

Complexity analysis

Provided the reducing function $f(n)$ is $\Theta(1)$

insert

$\Theta(1)$ reduction and $\Theta(1)$ memory operations

$$\Rightarrow \Theta(1)$$

find

$\Theta(1)$ reduction and $\Theta(1)$ memory operations

$$\Rightarrow \Theta(1)$$

delete

$\Theta(1)$ reduction and $\Theta(1)$ memory operations

$$\Rightarrow \Theta(1)$$

So what am I not telling you?

Sets of arbitrary things

- compute an integer (a *hash code*) for the things using a *hash function*
 - equal things **must** have equal hash codes
 - unequal things should be unlikely to share hash codes

computing an integer for the things:

Java `public int hashCode()`

C++ `operator()` functor second template argument to container

equal things must have equal integer codes:

Java `public boolean equals(Object o)`

C++ `operator()` functor third template argument to container

Work

1. Reading

- CLRS, sections 11.1 and 11.2
- DPV, section 1.5
- Drozdek, sections 10.1