

# Lecture 16

## Algorithms & Data Structures

Goldsmiths Computing

February 25, 2019

# Outline

Introduction

Fixed point

Multiplication

Floating point

# Outline

## Introduction

## Fixed point

## Multiplication

## Floating point

# Lecture

1. Random number generation
2. Comparison sorts
  - insertion sort
  - merge sort
  - heap sort
  - quick sort
  - $\Theta(N \log N)$  complexity bound
3. Shuffling
  - several broken shuffle versions
  - one correct  $\Theta(N \log N)$  shuffle
  - $\Theta(N)$  Fisher-Yates shuffle

# Lab

Catchup and future plans

# VLE activities

## Graphs quiz

# VLE activities (cont'd)

## Graphs quiz

# VLE activities (cont'd)

## Graphs quiz



# VLE activities (cont'd)

## Graphs submission

# Outline

Introduction

Fixed point

Multiplication

Floating point

# Motivation

Representing:

- integers within a range
- continuous sequence of place-value digits

Useful and practical:

- simple to implement in hardware
- reasonable properties
- ... but some unexpected behaviours too

# Definition

- data type for a number
- fixed number of digits
- implicit scaling factor
  - number represented is integer value multiplied by scaling factor

# Unsigned integers

## Fixed-width of binary digits

- assumed infinite leftward zeros
- radix (decimal) point at the right-hand end of the fixed-width field
- scaling factor of 1

## Representation

... 0 0 0 

0	0	0	0	0	0	1	1	0	1	0	0	1	0	0	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

(16-bit fixed point integer)

# Overflow

What happens when a calculation would put a 1 in the infinite sea of zeros?

**hardware** wraparound, carry flag

# Overflow

What happens when a calculation would put a 1 in the infinite sea of zeros?

**hardware** wraparound, carry flag

**C++** wraparound

# Overflow

What happens when a calculation would put a 1 in the infinite sea of zeros?

**hardware** wraparound, carry flag

**C++** wraparound

**Java** bad luck, no unsigned integers



# Signed integers

Many possible representations:

- sign-magnitude

**sign** high bit (0: positive; 1: negative)  
**magnitude** remaining bits

- ones complement

$$-x \quad \neg x$$

- twos-complement

$$-x \quad \neg x + 1$$

# Twos complement

What's so good about twos-complement?

- only one zero
- addition, subtraction and multiplication all the same as unsigned

## Representation

... 0 0 0 0 **0** 0 0 0 0 1 0 1 0 0 1 1 0 1 1 1

... 1 1 1 **1** 0 0 1 1 1 1 0 1 1 1 0 1 1 1 0

In practice, all current systems use twos-complement.

**function** NEG(x)

$r \leftarrow \neg x$

**return**  $r + 1$

**end function**

# Overflow

What happens when a calculation would put a non-sign bit in the infinite sea of sign bits?

**hardware** wraparound, carry flag

# Overflow

What happens when a calculation would put a non-sign bit in the infinite sea of sign bits?

**hardware** wraparound, carry flag

**Java** wraparound

# Overflow

What happens when a calculation would put a non-sign bit in the infinite sea of sign bits?

**hardware** wraparound, carry flag

**Java** wraparound

**C++** bad luck, undefined behaviour

# Absolute value

abs (C++ cstdlib) or Math.abs (Java):

- return (as an int) the absolute (non-negative) value of its argument

```
function ABS(x)
  if x < 0 then
    return NEG(x)
  else
    return x
  end if
end function
```

- does this always return a non-negative answer?

## Population count

popcnt (x86 instruction)

- return how many one bits are set in the (unsigned) integer argument.

Divide-and-conquer implementation:

```
function POPCNT(x)
```

```
    return POPCNTW(x,W)
```

▷ W is the integer width

```
end function
```

```
function POPCNTW(x,w)
```

```
    if w = 1 then
```

```
        return x
```

```
    else
```

```
        nw ←  $\frac{w}{2}$ 
```

▷ w assumed to be a power of 2

```
        lo ← POPCNTW(x &  $2^{nw} - 1$ , nw)
```

```
        hi ← POPCNTW( $\lfloor \frac{x}{2^{nw}} \rfloor$ , nw)
```

```
        return lo + hi
```

```
    end if
```

```
end function
```

# Population count

Parallel divide-and-conquer implementation:

**function** POPCNT(x)

▷ Assume W = 32

$x \leftarrow (x \& 0x55555555) + (\lfloor \frac{x}{2} \rfloor \& 0x55555555)$

$x \leftarrow (x \& 0x33333333) + (\lfloor \frac{x}{4} \rfloor \& 0x33333333)$

$x \leftarrow (x \& 0x0f0f0f0f) + (\lfloor \frac{x}{16} \rfloor \& 0x0f0f0f0f)$

$x \leftarrow (x \& 0x00ff00ff) + (\lfloor \frac{x}{256} \rfloor \& 0x00ff00ff)$

$x \leftarrow (x \& 0x0000ffff) + (\lfloor \frac{x}{65536} \rfloor \& 0x0000ffff)$

**return** x

**end function**



# Work

## 1. Reading

- Henry S. Warren, Jr. *Hacker's Delight*, Addison-Wesley (2003)
  - sections 5-1, 7-1

# Outline

Introduction

Fixed point

**Multiplication**

Floating point

# Motivation

- working with numbers as a data structure
- everyone knows how to multiply
- almost no-one knows how to multiply efficiently

## Previously

Numbers as array of digits (binary: bits)

- numbers have a *width*  $w$ , at least  $1 + \log_b(n)$

## Logical operations

`and(x,y)` return the bitwise logical and of  $x$  and  $y$

`xor(x,y)` return the bitwise exclusive-or of  $x$  and  $y$

## Previously

Numbers as array of digits (binary: bits)

- numbers have a *width*  $w$ , at least  $1 + \log_b(n)$

## Logical operations

`and(x,y)` return the bitwise logical and of  $x$  and  $y$

`xor(x,y)` return the bitwise exclusive-or of  $x$  and  $y$

## Arithmetic operations

`add(x,y)` return the sum of  $x$  and  $y$

`sub(x,y)` return the difference between  $x$  and  $y$

`shift(x,n)` return  $x$  multiplied by the base  $n$  times

## Previously

Numbers as array of digits (binary: bits)

- numbers have a *width*  $w$ , at least  $1 + \log_b(n)$

## Logical operations

`and(x,y)` return the bitwise logical and of  $x$  and  $y$

`xor(x,y)` return the bitwise exclusive-or of  $x$  and  $y$

## Arithmetic operations

`add(x,y)` return the sum of  $x$  and  $y$

`sub(x,y)` return the difference between  $x$  and  $y$

`shift(x,n)` return  $x$  multiplied by the base  $n$  times

## Complexity

- until now,  $\Theta(1)$
- in fact,  $\Theta(w) \sim \Theta(\log(n))$

(logarithmic factor is usually irrelevant, or width is taken as constant)

# Problem

Given these basic operations:

- how do we implement multiplication?
- how efficient is it?

# Example

123×135:



## Example

123×135:

## Primary (old-)school multiplication

			1	2	3
×			1	3	5
				1	5
		1	0		
		5			
				9	
		6			
	3				
		3			
	2				
	1				
	1	6	6	0	5

# Example

123×135:

Primary school multiplication

×	100	20	3
100	10000	2000	300
30	3000	600	90
5	500	100	15

# Complexity analysis

## time

For each digit in  $x$

- multiply with each digit in  $y$ .

Assume  $x$  and  $y$  are each of width  $w$

$$\Rightarrow \Theta(w^2)$$

# Divide and conquer

Write

- $x = x_{hi} \times b^{w/2} + x_{lo}$
- $y = y_{hi} \times b^{w/2} + y_{lo}$

Then  $x \times y = x_{hi}y_{hi} b^w + (x_{hi}y_{lo} + x_{lo}y_{hi})b^{w/2} + x_{lo}y_{lo}$

# Example


$$x \times y = x_{hi}y_{hi} b^w + (x_{hi}y_{lo} + x_{lo}y_{hi})b^{w/2} + x_{lo}y_{lo}$$

$$\begin{aligned} 123 \times 135 &= (1 \times 100 + 23) \times (1 \times 100 + 35) \\ &= 1 \times 1 \times 10000 + (23 \times 1 + 1 \times 35) \times 100 + 23 \times 35 \\ &= 10000 + 5800 + 805 \end{aligned}$$


# Recursion tree

$123 \times 135$

# Recursion tree




# Recursion tree







# Recursion tree



# Recursion tree



# Recursion tree



# Complexity analysis

time

$$x \times y = x_{hi}y_{hi} b^w + (x_{hi}y_{lo} + x_{lo}y_{hi})b^{w/2} + x_{lo}y_{lo}$$

- $T_w = 4T_{w/2} + O(w)$

# Complexity analysis

## time

$$x \times y = x_{hi}y_{hi} b^w + (x_{hi}y_{lo} + x_{lo}y_{hi})b^{w/2} + x_{lo}y_{lo}$$

- $T_w = 4T_{w/2} + O(w)$

## solution

Use the master theorem:

- $a = 4, b = 2, c = 1 \Rightarrow$  **case 1**
- $\log_b a = \log_2(4) = 2$

$$\Rightarrow T(w) \in \Theta(w^2)$$

# Divide and conquer, Karatsuba

Still with

- $x = x_{hi} \times b^{w/2} + x_{lo}$
- $y = y_{hi} \times b^{w/2} + y_{lo}$

Calculate

- $z_{hi} = x_{hi} \times y_{hi}$
- $z_{lo} = x_{lo} \times y_{lo}$
- $z_c = (x_{hi} + x_{lo}) \times (y_{hi} + y_{lo})$

Then  $x \times y = z_{hi} b^w + (z_c - z_{hi} - z_{lo})b^{w/2} + z_{lo}$

# Divide and conquer, Karatsuba

Still with

- $x = x_{hi} \times b^{w/2} + x_{lo}$
- $y = y_{hi} \times b^{w/2} + y_{lo}$

Calculate

- $z_{hi} = x_{hi} \times y_{hi}$
- $z_{lo} = x_{lo} \times y_{lo}$
- $z_c = (x_{hi} + x_{lo}) \times (y_{hi} + y_{lo})$

Then  $x \times y = z_{hi} b^w + (z_c - z_{hi} - z_{lo})b^{w/2} + z_{lo}$

## Example

123×135:

- $z_{hi} = 1$ ;  $z_{lo} = 35 \times 23 = 805$ ;  $z_c = 36 \times 24 = 864$

so:


$$\begin{aligned}
 123 \times 135 &= 1 \times 10000 + (864 - 805 - 1) \times 100 + 805 \\
 &= 10000 + 5800 + 805 \\
 &= 16605
 \end{aligned}$$

# Recursion tree


123×135




# Recursion tree




# Recursion tree




# Recursion tree



# Recursion tree



# Recursion tree



# Complexity analysis

time

$$x \times y = z_{hi} b^w + (z_c - z_{hi} - z_{lo})b^{w/2} + z_{lo}$$

- $T_w = 3T_{w/2} + O(w)$

# Complexity analysis

## time

$$x \times y = z_{hi} b^w + (z_c - z_{hi} - z_{lo})b^{w/2} + z_{lo}$$

- $T_w = 3T_{w/2} + O(w)$

## solution

Use the master theorem:

- $a = 3, b = 2, c = 1 \Rightarrow \text{case 1}$

- $\log_b a = \log_2(3) = 1.58$

$$\Rightarrow T(w) \in \Theta(w^{1.58})$$

# Work

## 1. Reading

- CLRS, section 4.2: matrix multiplication



# Outline

Introduction

Fixed point

Multiplication

Floating point

# Motivation

Represent a wider range of numbers than with fixed point

- fixed point: constant **absolute** precision
- floating point: constant **relative** precision

Convenient approximation to Real numbers

- but **only** an approximation
- ... some unexpected behaviours too

# Definition

Floating point behaviour is defined by an engineering standard:

- IEEE 754 (1985, revised 2008)

Implemented by most hardware platforms:

- floating-point units in CPUs
  - (software support for FPU features varies)
- graphics cards (CUDA, OpenGL)
  - and other coprocessors

# Operations

## Operations on floats:

- add** return the sum of two floating point numbers
- sub** return the difference of two floating point numbers
- mul** return the product of two floating point numbers
- div** return the quotient of two floating point numbers
- sqrt** return the square root of one floating point number

## Operations on floating point units:

- rounding mode** should rounding go towards  $+\infty$ , 0,  $-\infty$  or even?
- trapping** should the FPU generate an exception for conditions such as overflow or divide by zero?

# General idea

Represent a number  $n$  as:

$$n = \text{sign} \times \text{significand} \times 2^{\text{exponent}}$$

**sign** 1 or -1

**significand** number in  $[1, 2)$

**exponent**  $\left\lfloor \log_2 n \right\rfloor$

Represent this in a fixed-size field using:

**sign bit** 0 (positive) or 1 (negative)

**mantissa** *fractional* part of significand

**exponent** exponent + bias

$$n = (-1)^s \times (1 + m) \times 2^{e-B}$$

# Single-precision

- 32-bit quantity:
  - 1 sign bit
  - 8 exponent bits
    - bias is 127, range is  $\pm 2^{-126}$  to  $2^{127}$
  - 23 mantissa bits
    - plus “hidden bit” gives 24 binary (~7 decimal) digits of precision



# Single-precision

- 32-bit quantity:
  - 1 sign bit
  - 8 exponent bits
    - bias is 127, range is  $\pm 2^{-126}$  to  $2^{127}$
  - 23 mantissa bits
    - plus “hidden bit” gives 24 binary (~7 decimal) digits of precision

## Representation



## Example

$$0.5 = 1 \times (1 + 0) \times 2^{-1}$$

sign 0

mantissa 0

exponent 126 (0x7e)

overall 0x3f000000



# Zero?

No representation for zero in this scheme

$$\left\lfloor \log_2(x) \right\rfloor = -\infty$$

Special representation of zero:

exponent field 0

mantissa 0

sign 0 or 1

# Double-precision

- 64-bit quantity:
  - 1 sign bit
  - 11 exponent bits
    - bias is 1023, range is  $\pm 2^{-1022}$  to  $2^{1023}$
  - 52 mantissa bits
    - plus “hidden bit” gives 53 binary (~16 decimal) digits of precision

# Double-precision

- 64-bit quantity:
  - 1 sign bit
  - 11 exponent bits
    - bias is 1023, range is  $\pm 2^{-1022}$  to  $2^{1023}$
  - 52 mantissa bits
    - plus “hidden bit” gives 53 binary (~16 decimal) digits of precision

## Representation



# Double-precision

- 64-bit quantity:
  - 1 sign bit
  - 11 exponent bits
    - bias is 1023, range is  $\pm 2^{-1022}$  to  $2^{1023}$
  - 52 mantissa bits
    - plus “hidden bit” gives 53 binary (~16 decimal) digits of precision

## Representation



## Example

$$0.75 = 1 \times (1 + 0.5) \times 2^{-1}$$

sign 0

mantissa shift(1, 51)

exponent 1022 (0x3fe)

overall 0x3fe8000000000000

# Epsilon

Floating point has a larger *range* than *precision*

- calculations with floating points will usually not give an exactly representable answer
  - (even if the input numbers were exact)

## Epsilon

$\epsilon$  is the smallest float which you can add to 1.0 and get an answer that isn't 1.0:

single-precision  $2^{-24} + 2^{-47}$

double-precision  $2^{-53} + 2^{-105}$

For all  $0 < x < \epsilon$

$$1 + x \rightarrow 1$$

# Inverse square root

## Game programming history

- need to take  $f(x) = \frac{1}{\sqrt{x}}$  often and quickly

Use  $\log_2((1 + m) \times 2^{e-B}) \approx e - B + m$ :

```
int i = *(int *) &f;
i = 0x5f3759df - (i >> 1);
f = *(float *) &i;
```

This was good in 1999 (*Quake III Arena*)

- nowadays we have hardware to do this
- SSE rsqrtss

# Work

## 1. Reading

- David Goldberg, *What every computer scientist should know about floating point arithmetic*, Computing (1991)